



# Bayesian networks (BNs)

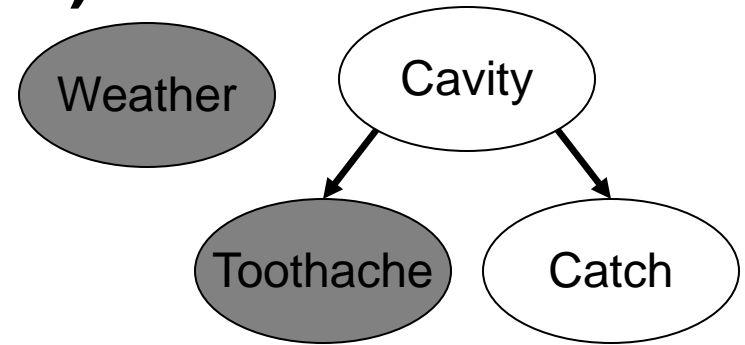
- A type of *graphical model*
- A BN states conditional independence relationships between random variables
- Compact specification of full joint distributions



# Syntax

- **Directed Acyclic Graph (DAG)**

- **Nodes:** random variables
  - Can be assigned (observed) or unassigned (unobserved)




- **Arcs:** interactions

- An arrow from one variable to another indicates direct influence
- Encode conditional independence
  - *Weather* is independent of the other variables
  - *Toothache* and *Catch* are conditionally independent given *Cavity*
- Must form a directed, *acyclic* graph


# Example: N independent coin flips

- Complete independence: no interactions



# Example: Naïve Bayes document model

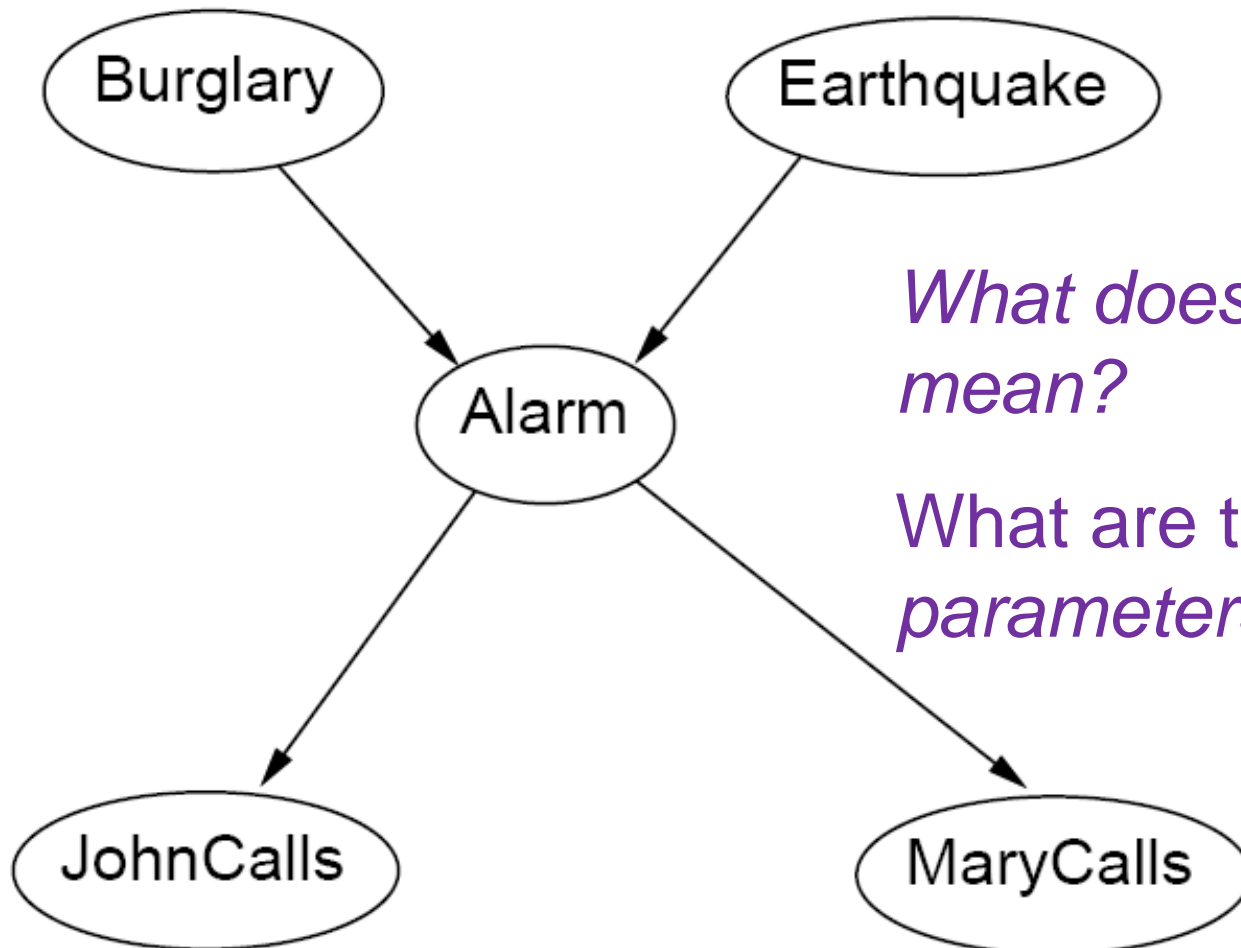
- Random variables:
  - $X$ : document class
  - $W_1, \dots, W_n$ : words in the document



# Example: Burglar Alarm

- I have a burglar alarm that is sometimes set off by minor earthquakes. My two neighbors, John and Mary, promised to call me at work if they hear the alarm
  - Example inference task: suppose Mary calls and John doesn't call. What is the probability of a burglary?
- What are the random variables?
  - *Burglary, Earthquake, Alarm, JohnCalls, MaryCalls*
- What are the direct influence relationships?
  - A burglar can set the alarm off
  - An earthquake can set the alarm off
  - The alarm can cause Mary to call
  - The alarm can cause John to call

# Example: Burglar Alarm




*What does this mean?*

*What are the model parameters?*

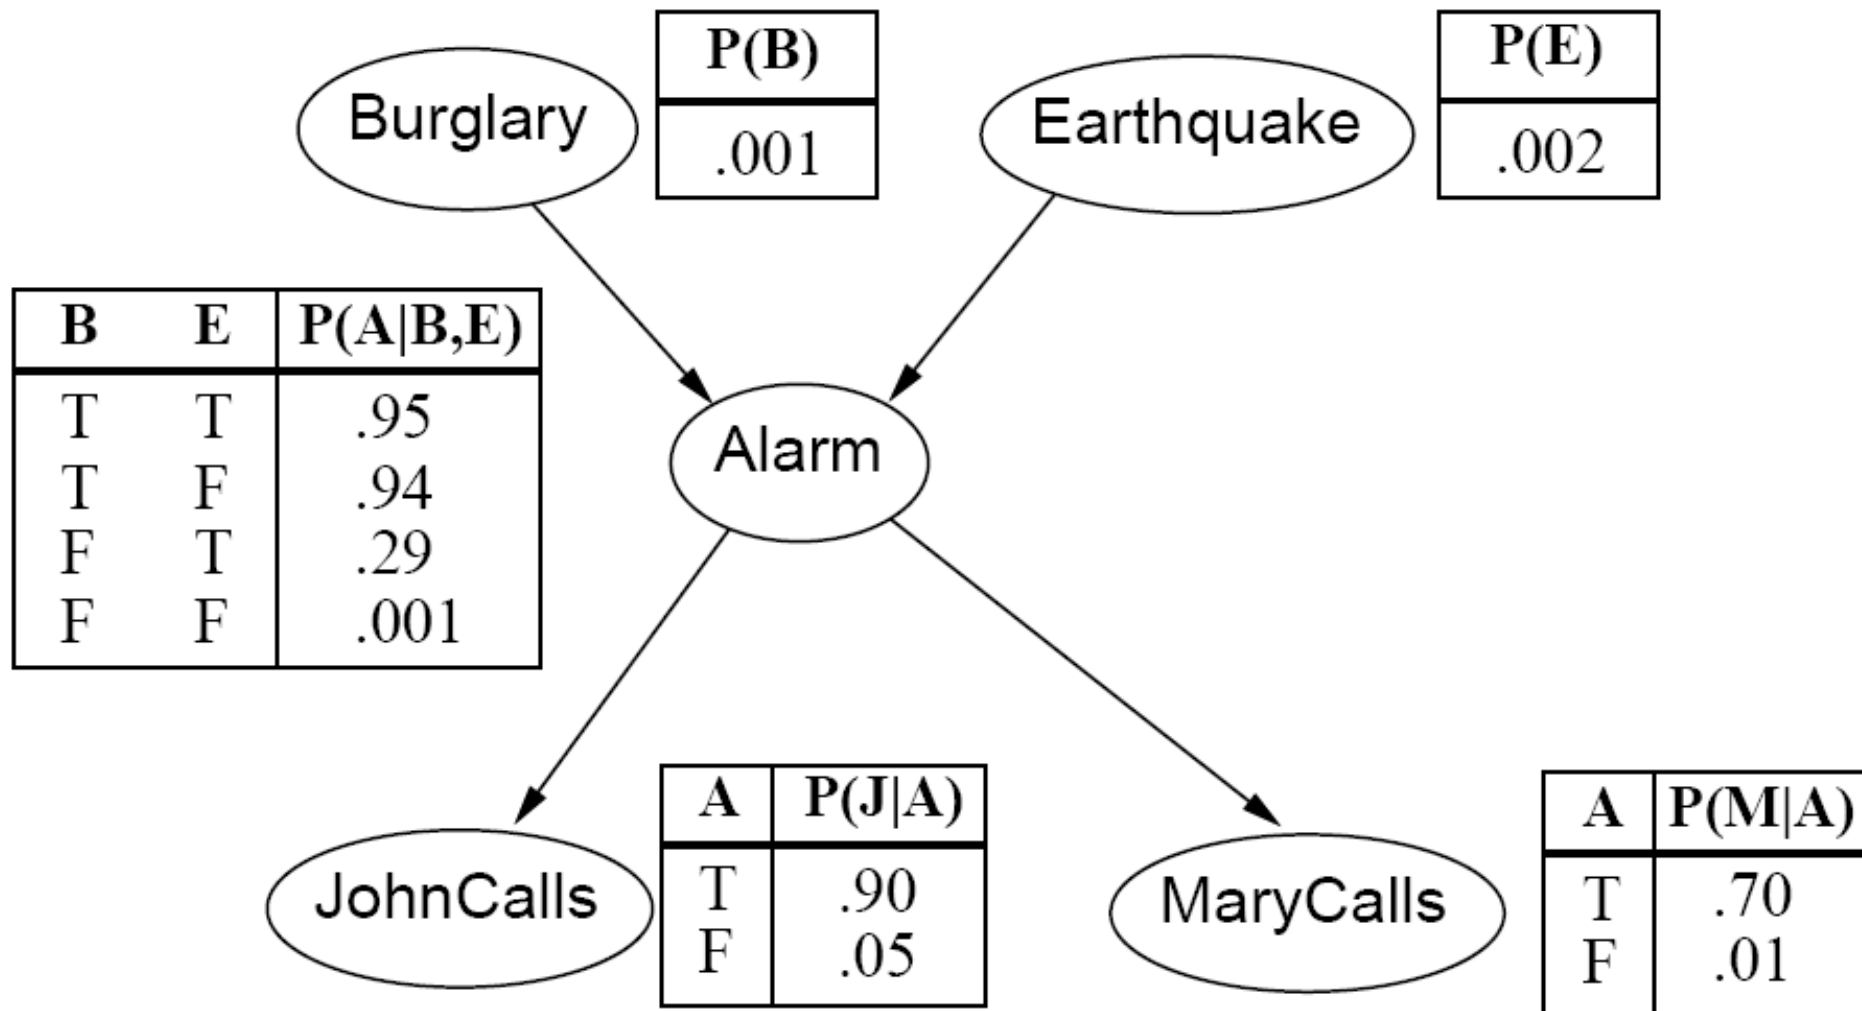
# Semantics

- A BN represents a full joint distribution in a compact way.
- We need to specify a *conditional probability* distribution for each node given its parents:

$$P(X \mid \text{Parents}(X))$$



# Example: Burglar Alarm





# The joint probability distribution


- Key property: each node is conditionally independent of its non-descendents given its parents
- Suppose the nodes  $X_1, \dots, X_n$  are sorted in topological order
- For each node  $X_i$ , we know  $P(X_i \mid \text{Parents}(X_i))$
- To get the joint distribution  $P(X_1, \dots, X_n)$ , use chain rule:

$$\begin{aligned}
 P(X_1, \dots, X_n) &= \prod_{i=1}^n P(X_i \mid X_1, \dots, X_{i-1}) \\
 &= \prod_{i=1}^n P(X_i \mid \text{Parents}(X_i))
 \end{aligned}$$

# The joint probability distribution


$$P(X_1, \dots, X_n) = \prod_{i=1}^n P(X_i \mid P \setminus X_i)$$

- For example,  $P(j, m, a, \neg b, \neg e)$
- $= P(\neg b) P(\neg e) P(a \mid \neg b, \neg e) P(j \mid a) P(m \mid a)$



# Conditional independence

- Key property:  $X$  is conditionally independent of every *non-descendant node* given its parents
- Causal Chain



X: Low pressure

Y: Rain


Z: Traffic

- Are  $X$  and  $Z$  independent?
- Is  $Z$  independent of  $X$  given  $Y$ ?

$$P(Z \mid X, Y) = \frac{P(X, Y, Z)}{P(X, Y)} = \frac{P(X)P(Y \mid X)P(Z \mid Y)}{P(X)P(Y \mid X)} = P(Z \mid Y)$$

# Conditional independence

Common cause



Y: Project due

X: Newsgroup  
busy

Z: Lab full


Are X and Z independent?

No

Are they conditionally  
independent given Y?

Yes

Common effect



X: Raining

Z: Ballgame

Y: Traffic


Are X and Z independent?

Yes

Are they conditionally  
independent given Y?

No

# Compactness



- Suppose we have a Boolean variable  $X_i$  with  $k$  Boolean parents. How many rows does its conditional probability table have?
  - $2^k$  rows for all the combinations of parent values
  - Each row requires one number  $p$  for  $X_i = \text{true}$
- If each variable has no more than  $k$  parents, how many numbers does the complete network require?
  - $O(n \cdot 2^k)$  numbers – vs.  $O(2^n)$  for the full joint distribution
- How many nodes for the burglary network?  
 $1 + 1 + 4 + 2 + 2 = 10$  numbers (vs.  $2^5 - 1 = 31$ )

# Constructing Bayesian networks

1. Choose an ordering of variables  $X_1, \dots, X_n$
2. For  $i = 1$  to  $n$ 
  - add  $X_i$  to the network
  - select parents from  $X_1, \dots, X_{i-1}$  such that
$$P(X_i \mid \text{Parents}(X_i)) = P(X_i \mid X_1, \dots, X_{i-1})$$

# Example

- Suppose we choose the ordering M, J, A, B, E

$$P(J \mid M) = P(J)?$$

MaryCalls


JohnCalls

# Example

- Suppose we choose the ordering M, J, A, B, E

$$P(J \mid M) = P(J)?$$

No





# Example

- Suppose we choose the ordering M, J, A, B, E


$$P(J \mid M) = P(J)?$$

No

$$P(A \mid J, M) = P(A)?$$

$$P(A \mid J, M) = P(A \mid J)?$$

$$P(A \mid J, M) = P(A \mid M)?$$



# Example

- Suppose we choose the ordering M, J, A, B, E

$$P(J \mid M) = P(J)?$$

No

$$P(A \mid J, M) = P(A)?$$


No

$$P(A \mid J, M) = P(A \mid J)?$$

No

$$P(A \mid J, M) = P(A \mid M)?$$

No



# Example

- Suppose we choose the ordering M, J, A, B, E

$$P(J \mid M) = P(J)?$$

No

$$P(A \mid J, M) = P(A)?$$

No

$$P(A \mid J, M) = P(A \mid J)?$$


No

$$P(A \mid J, M) = P(A \mid M)?$$

No

$$P(B \mid A, J, M) = P(B)?$$

$$P(B \mid A, J, M) = P(B \mid A)?$$



# Example

- Suppose we choose the ordering M, J, A, B, E

$$P(J \mid M) = P(J)?$$

No

$$P(A \mid J, M) = P(A)?$$

No

$$P(A \mid J, M) = P(A \mid J)?$$

No

$$P(A \mid J, M) = P(A \mid M)?$$


No

$$P(B \mid A, J, M) = P(B)?$$

No

$$P(B \mid A, J, M) = P(B \mid A)?$$

Yes



# Example

- Suppose we choose the ordering  $M, J, A, B, E$

$$P(J \mid M) = P(J)?$$

No

$$P(A \mid J, M) = P(A)?$$

No

$$P(A \mid J, M) = P(A \mid J)?$$

No

$$P(A \mid J, M) = P(A \mid M)?$$

No

$$P(B \mid A, J, M) = P(B)?$$


No

$$P(B \mid A, J, M) = P(B \mid A)?$$

Yes

$$P(E \mid B, A, J, M) = P(E)?$$

$$P(E \mid B, A, J, M) = P(E \mid A, B)?$$



# Example

- Suppose we choose the ordering M, J, A, B, E

$$P(J \mid M) = P(J)?$$

No

$$P(A \mid J, M) = P(A)?$$

No

$$P(A \mid J, M) = P(A \mid J)?$$

No

$$P(A \mid J, M) = P(A \mid M)?$$

No

$$P(B \mid A, J, M) = P(B)?$$

No

$$P(B \mid A, J, M) = P(B \mid A)?$$


Yes

$$P(E \mid B, A, J, M) = P(E)?$$


No

$$P(E \mid B, A, J, M) = P(E \mid A, B)?$$

Yes



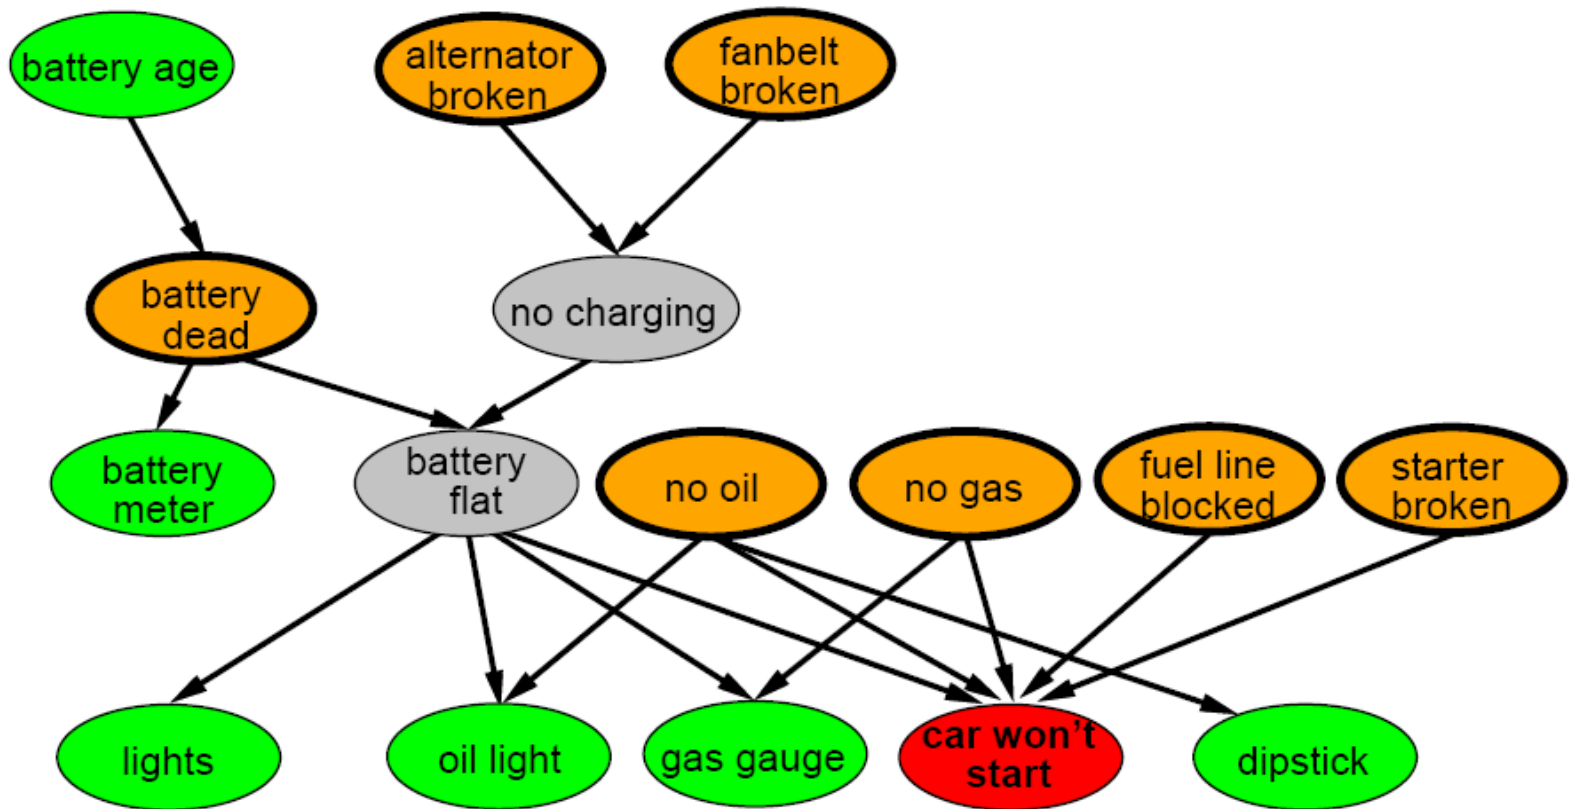
# Example contd.



- Deciding conditional independence is hard in noncausal directions
  - The causal direction seems much more natural, but is not mandatory
- Network is less compact:  $1 + 2 + 4 + 2 + 4 = 13$  numbers needed

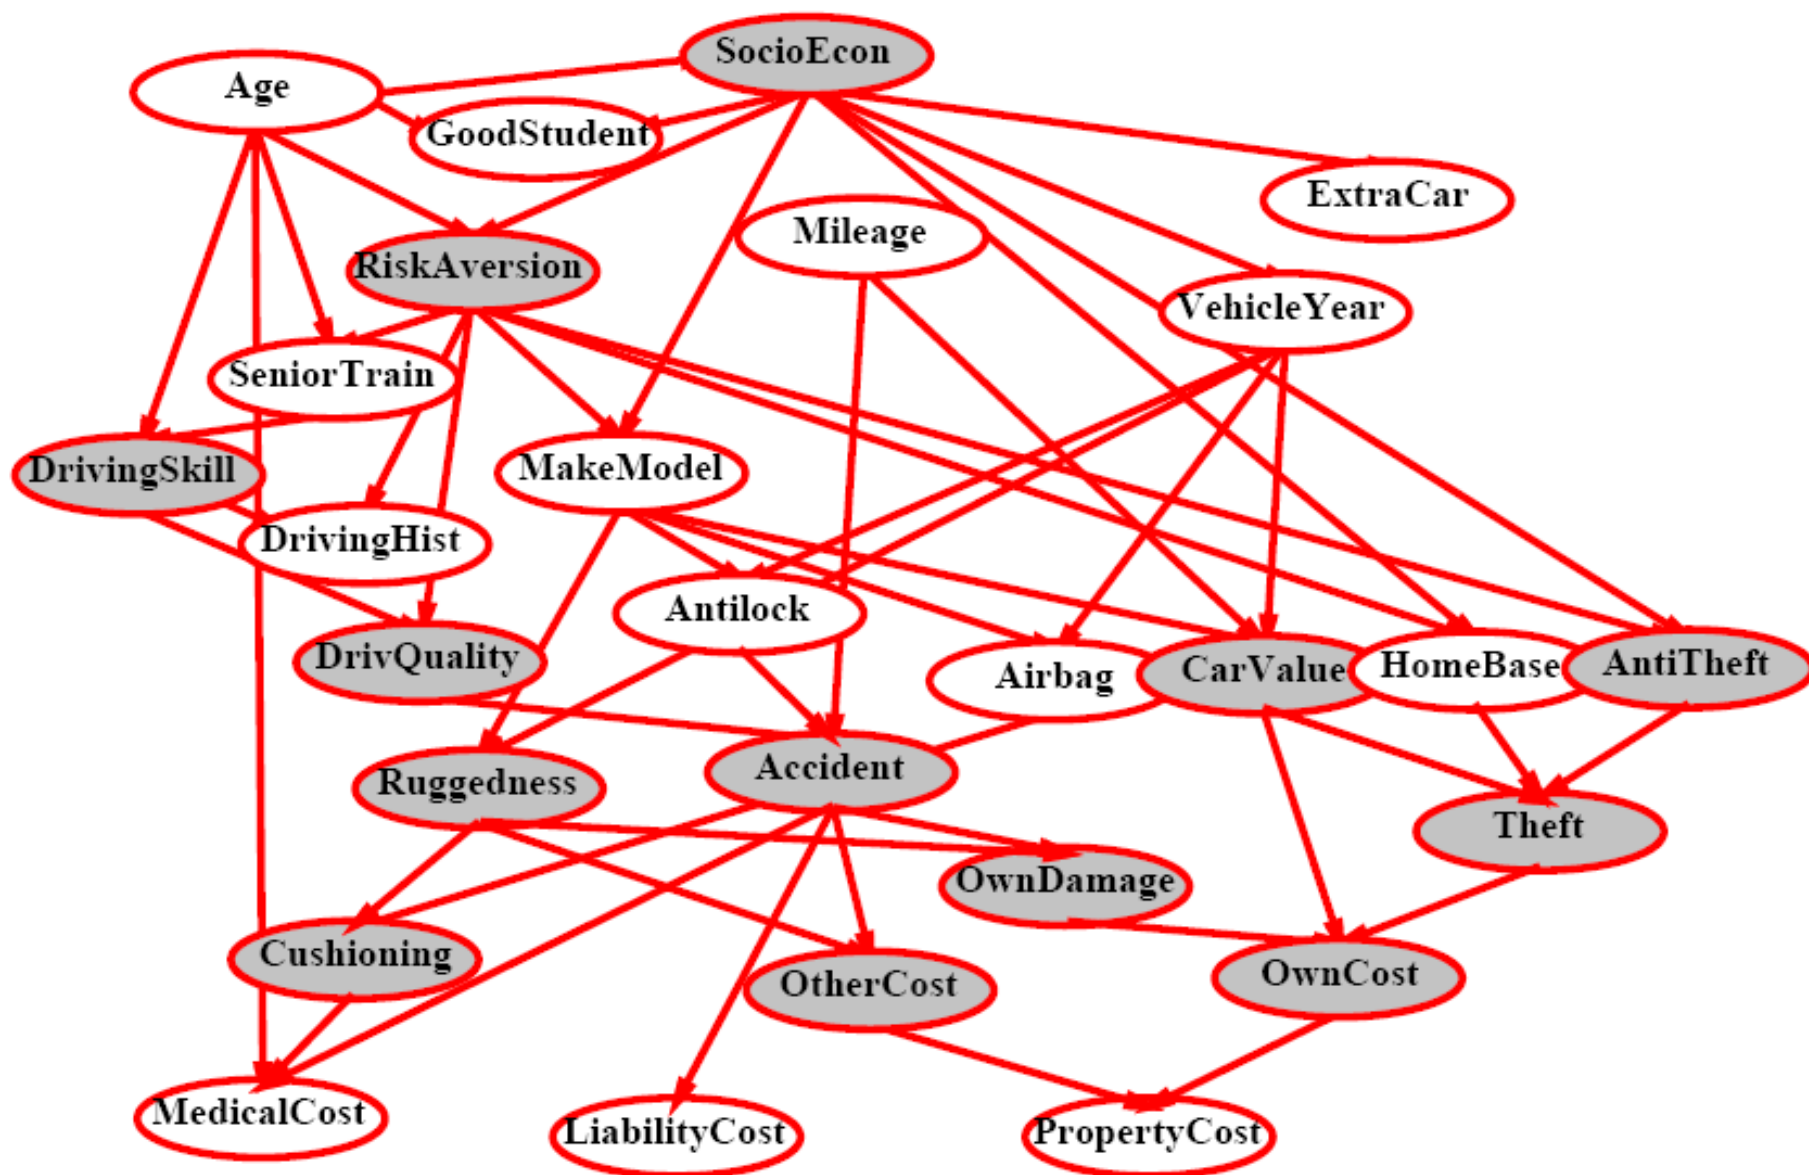
# A more realistic Bayes Network: Car diagnosis

- **Initial observation:** car won't start
- **Orange:** “broken, so fix it” nodes
- **Green:** testable evidence
- **Gray:** “hidden variables” to ensure sparse structure, reduce parameters






# Car insurance



# In computational biology...



## Causal Protein-Signaling Networks Derived from Multiparameter Single-Cell Data

Karen Sachs, Omar Perez, Dana Pe'er, Douglas A. Lauffenburger, and Garry P. Nolan  
(22 April 2005) *Science* **308** (5721), 523.


# Identifying Aggressive Behavior




World Phenomena  
(e.g., bullying)



Game




Inference



Inferred Roles

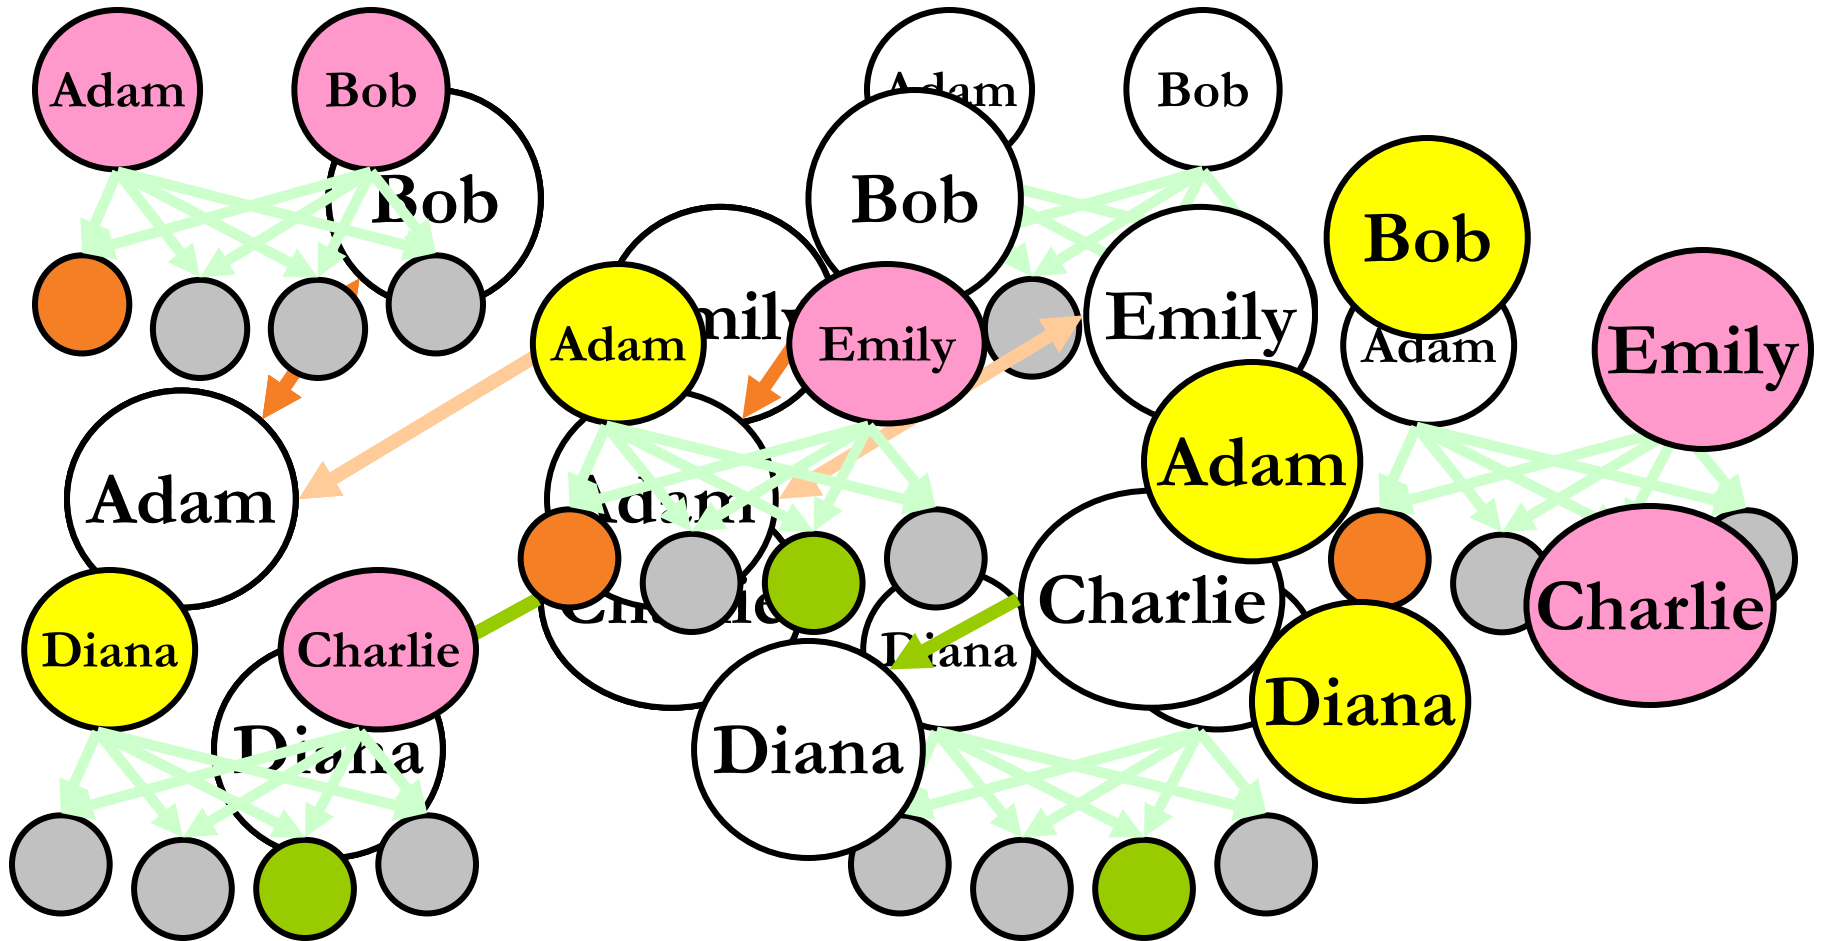
Mancilla-Caceres, J.F., Pu, W., Amir, E., and D. Espelage. *Identifying Bullies with a Computer Game*. In Proceedings of the 26th AAAI Conference on Artificial Intelligence (AAAI-12). 2012

# Pairwise Model for Identification of Aggressive Behavior

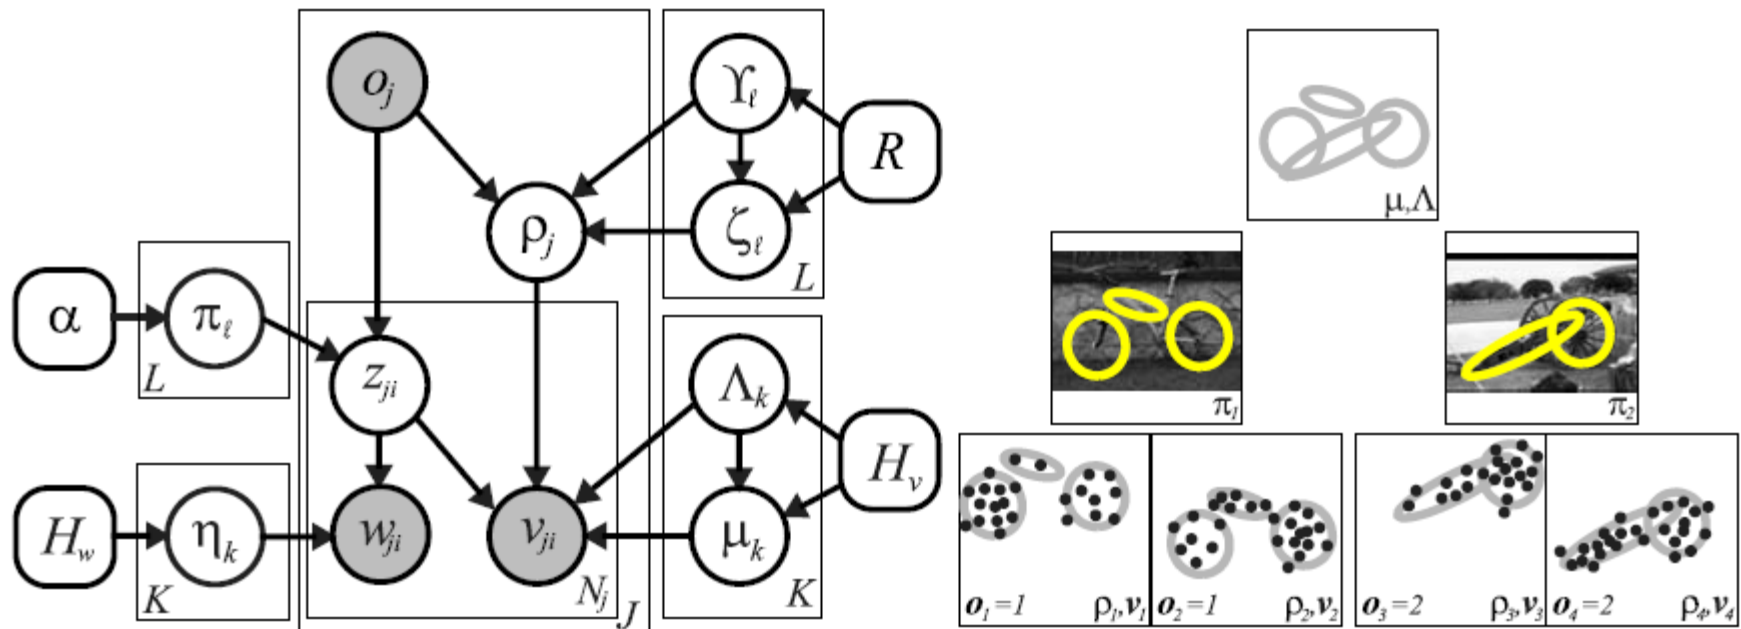


- **Wins** = 1 if *Adam* got more coins from *Bob* than *Bob* from *Adam*.
- **Coercive** = 1 if *Adam* sent more coercive than prosocial messages to *Bob*.
- **Negative Affect** = 1 if *Adam* sent more negative affect messages than positive ones to *Bob*.
- **Nominated** = 1 if *Adam* nominated *Bob* to be on the same team, -1 if the nomination was negative, 0 otherwise.
- **Same Team** = 1 if *Adam* and *Bob* belong to the same team.

# Global Inference from Pairwise Interactions.



# In computer vision



**Fig. 3** A parametric, fixed-order model which describes the visual appearance of  $L$  object categories via a common set of  $K$  shared parts. The  $j^{th}$  image depicts an instance of object category  $o_j$ , whose position is determined by the reference transformation  $\rho_j$ . The appearance  $w_{ji}$  and position  $v_{ji}$ , relative to  $\rho_j$ , of visual features are determined


by assignments  $z_{ji} \sim \pi_{o_j}$  to latent parts. The cartoon example illustrates how a wheel part might be shared among two categories, *bicycle* and *cannon*. We show feature positions (but not appearance) for two hypothetical samples from each category

## **Describing Visual Scenes Using Transformed Objects and Parts**

E. Sudderth, A. Torralba, W. T. Freeman, and A. Willsky.

International Journal of Computer Vision, No. 1-3, May 2008, pp. 291-330.

# Continuous Variables Example



$$P(c|h, \text{subsidy}) = N(a_t h + b_t, \sigma_t^2)(c) = \frac{1}{\sigma_t \sqrt{2\pi}} e^{-\frac{1}{2} \left( \frac{c - (a_t h + b_t)}{\sigma_t} \right)^2}$$

$$P(c|h, \neg \text{subsidy}) = N(a_f h + b_f, \sigma_f^2)(c) = \frac{1}{\sigma_f \sqrt{2\pi}} e^{-\frac{1}{2} \left( \frac{c - (a_f h + b_f)}{\sigma_f} \right)^2}$$

$$P(\text{buys} \mid \text{Cost} = c) = \frac{1}{1 + \exp(-2 \frac{-c + \mu}{\sigma})}$$

# Summary

- Bayesian networks provide a natural representation for (causally induced) conditional independence
- Topology + conditional probability tables
- Generally easy for domain experts to construct



# Probabilistic inference

- A general scenario:
  - Query variables:  $\mathbf{X}$
  - Evidence (observed) variables:  $\mathbf{E} = \mathbf{e}$
  - Unobserved variables:  $\mathbf{Y}$
- If we know the full joint distribution  $P(\mathbf{X}, \mathbf{E}, \mathbf{Y})$ , how can we perform inference about  $\mathbf{X}$ ?

$$P(\mathbf{X} \mid \mathbf{E} = \mathbf{e}) = \frac{P(\mathbf{X}, \mathbf{e})}{P(\mathbf{e})} \propto \sum_{\mathbf{y}} P(\mathbf{X}, \mathbf{e}, \mathbf{y})$$

- Problems
  - Full joint distributions are too large
  - Marginalizing out  $\mathbf{Y}$  may involve too many summation terms